# AI Content, Batch 01

Six honest, source-backed content experiments built to earn real conversations, saves and shares, run as a learning loop instead of six random posts.

Six posts show direction, not proof. Every number in this document traces to a verified source.

| BATCH | DATE | CHANNELS | PRIMARY METRICS |
|---|---|---|---|
| 01 | 2026-08-05, Riyadh | X, Instagram, LinkedIn | Saves, shares, substantive replies, real conversations |

# 01 What this is

Six content experiments, each grounded in a real, fetchable, current source, each built to invite a genuine opinion rather than manufacture a fight.

Muse's marketing is a pillar, not a task, because it earns nothing directly and would otherwise never happen. It exists for inbound client work, a launch audience for ventures, institutional access, and standing for the people here. Personal founder and team accounts do the heavy lifting. The Muse account amplifies and holds the official archive. Nobody is required to post.

## The tension, named plainly

The goal is real reach and engagement. Muse also rejects engagement bait, comment pods, fake conflict, reposting other people's work as its own, and empty farming. Those two things read as opposites until you notice the actual lever: honest, source-backed posts that ask for a real opinion or a useful disagreement earn attention without the cheap tactics. Trust is the asset the whole pillar depends on, so it is not spent for a single post's reach.

> **WHERE AI SITS IN THIS**
>
> Muse uses AI where it changes the outcome and never claims it. This batch is neither anti-AI nor hype. Two experiments argue AI is genuinely useful right now. Two are critical of a specific AI claim or trend. Two are AI research findings taken on their own terms. A slate that read as uniformly cynical about AI would itself be a brand mismatch.

The number that matters is real conversations, then saves and shares, then repeat viewers, then audience composition, then consistency. Followers, likes, impressions and reach are tracked and never treated as success on their own.

## The five content pillars this batch draws from

| Useful craft | Working in the open | Opinion, stated fairly |
|---|---|---|
| A picking heuristic, a testing habit, a workflow worth trying this week | Naming the counter-argument, the conflict of interest, the thing that is still unproven | A defensible position with its own rebuttal attached, built to be argued with rather than merely liked |

Free assets and behind-the-scenes posting are the other two pillars in Muse's engine, not exercised directly by this batch but part of the same rotation over time.

# 02 Strategy, and how a topic earns a slot

A candidate had to clear every selection rule below. Any one rejection rule was enough to kill it.

## Selection

- [ ] Every claim traces to a real, fetchable, current source with a verbatim excerpt
- [ ] It invites a genuine opinion or lived experience, not manufactured conflict
- [ ] It sits in Muse's territory: design and code craft, AI quality and measurement, Arabic and RTL, product quality
- [ ] Practitioner-first relevance for designers, developers, product people, AI builders, students and GCC studios
- [ ] Format-ready as a 4 to 5 slide carousel, an X thread, or both
- [ ] Honest register: we can say where the thing is weak, not only where it shines

## Rejection

- **Unfetchable or unconfirmable.** The source could not be fetched, or a number, quote or date could not be confirmed from the live page
- **Vendor spin as fact.** Or a consequential single-source claim with no independent corroborator
- **Bait.** Reads as engagement bait, outrage without substance, or a brand advertisement
- **Named attack.** Attacks a named Saudi company, or risks reading that way
- **Borrowed insight.** Reposts someone else's finding as ours, or needs a banned phrase to work

> The batch spans more than one content pillar on purpose: useful craft, free assets, behind the scenes, working in the open, and opinion or market commentary.

| RULE | WHAT IT PROTECTS AGAINST |
| --- | --- |
| **Re-verify at build time** | Plan.md is reconnaissance, not evidence. Pages change, so every claim used in final copy was re-fetched live on 2026-08-05. |
| **Cut, do not guess** | If a page cannot be fetched or a number cannot be confirmed, the claim is cut, not quoted from memory. |

# 03 The slate at a glance

Six experiments, spanning current AI news, a usable workflow, a research finding, behavioural science, a defensible opinion, and Muse's own Arabic and RTL territory.

| # | TOPIC | FORMAT | PRIMARY METRIC | SOURCE AGE |
|---|---|---|---|---|
| 1 | Three model launches, picking without the leaderboard | X and Instagram carousel | **Saves** | In 90 days |
| 2 | Agent worktrees, parallel AI coding | X thread, carousel | **Saves** | In 90 days |
| 3 | ReasonBENCH, benchmark score instability | Instagram carousel, X post | **Substantive replies** | In 90 days |
| 4 | AI trust, appreciation and aversion | Instagram carousel, X thread | **Substantive replies** | Mixed, disclosed |
| 5 | AI code volume vs maintainability | X thread | **Quote posts** | Mixed, disclosed |
| 6 | Arabic benchmark quality, native models | Instagram carousel, X post | **Real conversations** | Slightly past 90d, disclosed |

## Mix and recency

Four of six experiments rest on a lead source inside the last 90 days: 1, 2, 3 and part of 5. Experiment 4 pairs a 2026 finding with two pre-LLM behavioural landmarks, labelled by year. Experiment 6 uses slightly older 2026 material. The mix spans model news, workflow, evaluation research, behavioural science, a defensible opinion, and Muse's Arabic territory.

Shared distribution model across all six: a named founder or team member posts first from a personal account. The Muse account amplifies within a few hours with a one line framing and keeps the archive. LinkedIn receives an adapted version a day or two later, once early replies have shaped the angle.

# 04 Experiment 1. Model picking

In one month, three labs shipped major model updates. Each gives you a different reason to switch: a new default, a benchmark lead, or lower-token Flash models.

> "Three major AI releases landed in one month. Each gives you a different reason to switch."

- **1.** Three major AI releases landed in one month. Each gives you a different reason to switch.
- **2.** Claude Opus 5 [1], GPT-5.6 Sol, Terra, Luna [2], Gemini 3.6 Flash, 3.5 Flash-Lite, 3.5 Flash Cyber [3]. Each launch sells a different win.
- **3.** Anthropic calls Opus 5 its new default on Claude Max [1]. OpenAI says Sol sets a new state of the art on its own benchmark [2]. Google says its Flash trio is not even the flagship line, that is Gemini Pro [3].
- **4.** Pick for your work, not the leaderboard: your task, your language including Arabic, your latency and cost, tested on your own examples.
- **5.** Which model actually holds up on your real work, and what do you route where?

**THE QUESTION**

Which model actually holds up on your real work, and what do you route where?

## Why it can engage

- **Comments.** A low-effort way to name an actual default model and what it routes to
- **Saves.** A picking heuristic that outlasts this month's news cycle
- **Shares.** Sent into a team argument about routing
- **Quote posts.** A reader's own selection rule, added on top

Primary metric: saves. Sources: Anthropic 2026-07-24 [1], TechCrunch 2026-07-09 [2] and 2026-07-21 [3]. Full register in sources/sources.md.

# 05 Experiment 2. Agent worktrees

A workflow available now, not theoretical: each AI coding session gets its own Git worktree, an isolated copy of the repo. It genuinely helps, and it has a real downside.

> ”You can now run several AI coding agents at once without them fighting over your repo.”

- **1.** Running two or three AI agents at once used to mean chaos in one working tree.
- **2.** New in GitHub Copilot, July 2026: start a Copilot, Claude or Codex session in its own Git worktree, an isolated copy of your repository [4].
- **3.** Why it helps: each agent works on its own copy and branch. Nothing touches main until you review and merge.
- **4.** Where it bites: one developer had two agents ”stomp on each other’s edits inside about ninety seconds” before worktrees existed [27]. Isolation is not the same as done.
- **5.** Do you run agents in parallel yet? What is your review-and-merge routine before it turns into branch soup?

> **THE QUESTION**
>
> How do you keep parallel agents from becoming a pile of half-finished branches?

## Why it can engage

- **Comments.** Real parallel-agent setups, and what broke before this existed
- **Saves.** A workflow worth trying on a Sunday
- **Shares.** Sent to a team lead already fielding agent merge conflicts
- **Quote posts.** A merge discipline of the reader’s own

Primary metric: saves. Sources: GitHub Changelog 2026-07-30 [4], corroborated by dev.to, Leo Baniak, 2026-07-31 [27]. Full register in sources/sources.md.

# 06 Experiment 3. Benchmark variance

A 2026 paper found reasoning-strategy scores unstable across repeated runs, even at zero temperature. A single leaderboard number can silently misrank two close systems.

> "Run the same benchmark twice and the winner can change."

- **1.** We treat one benchmark score as the answer. It is one sample.
- **2.** ReasonBENCH, an arXiv paper revised May 2026, ran 10 reasoning strategies across 12 models and 6 tasks, 30 independent trials each [6].
- **3.** The top strategy won only 77% of head-to-head runs against its nearest competitor [6]. That is a win rate across runs, not an accuracy score.
- **4.** Two close systems can swap places from one run to the next. Run your eval more than once, report a range, distrust a narrow win.
- **5.** Have you seen a model win one day and lose the next on the same test? What do you do about eval variance?

THE QUESTION

How do you handle it when the same eval ranks models differently on a re-run?

## Why it can engage

- **Comments.** Eval-flakiness stories from people who run tests for a living
- **Saves.** A testing rule of thumb worth keeping
- **Shares.** Sent to whoever picks the team's model
- **Quote posts.** Arguing whether leaderboards mean anything at all

Primary metric: substantive replies. Source: arXiv 2512.07795, v2 2026-05-30 [6]. Author affiliations could not be confirmed cleanly and are not named in copy. Full register in sources/sources.md.

# 07 Experiment 4. AI trust

A current workplace-AI finding and two behavioural landmarks reveal a useful tension: we can defer to algorithms too easily, then reject them too quickly after one visible mistake.

> "We trust algorithms too much. Then one mistake makes us trust them too little."

- **1.** We trust algorithms too much. Then one mistake makes us trust them too little.
- **2.** A 2026 study: 58% said AI "did most of the thinking." They reported less confidence and ownership. People who challenged AI reported stronger authorship [28].
- **3.** Algorithm appreciation, 2019: people followed advice more when they believed it came from an algorithm than from a person [23].
- **4.** Algorithm aversion, 2015: after the same mistake, people lost confidence in an algorithm faster than in a human, even when the algorithm performed better [22].
- **5.** Where do you switch from trusting AI to rejecting it, and is that line based on evidence or one bad answer?

> **THE QUESTION**
>
> Where do you switch from trusting AI to rejecting it, and what evidence moves that line?

## Why it can engage

- **Comments.** The AI mistake that changed someone's behaviour
- **Saves.** A memorable behavioural paradox worth carrying into the next project
- **Shares.** Sent to a teammate who either trusts too quickly or rejects after one failure
- **Quote posts.** A reader's own rule for verification and override

Primary metric: substantive replies. Sources: Tech Xplore 2026-04-16 [28], Dietvorst, Simmons and Massey 2015 [22], Logg, Minson and Moore 2019 [23]. The older studies are pre-LLM behavioural concepts, not chatbot evaluations. Full register in sources/sources.md.

# 08 Experiment 5. Code maintainability

Muse's own opinion, stated fairly: AI is raising code volume and speed while eroding the signals that mark maintainable code, and we are poor judges of our own speed while it happens.

> "AI is helping us ship more code, faster. The maintainability is going the other way."

- **1.** AI is helping us write more code, faster. That part is real. The quality signals are the worry.
- **2.** Reported by LeadDev from GitClear research: code duplication up 81%, refactoring down 70%, legacy refactoring down 74% since 2023 [8]. GitClear sells code-quality tooling, a real conflict of interest, stated plainly.
- **3.** A METR trial: experienced developers took 19% longer with AI, yet still believed they were 20% faster [9]. METR later redesigned its continuing study after newer recruitment selection effects made current estimates unreliable [10].
- **4.** The honest frame, DORA 2025: "AI is an amplifier, not a fix" [11]. It rewards strong foundations and punishes weak ones.
- **5.** Counter-argument, fairly: METR's sample was small and self-selected, and GitClear sells the tool that measures this problem. Six months on, is your AI-assisted code easier or harder to maintain than what you wrote by hand?

> **THE QUESTION**
>
> Six months on, is your AI-assisted code easier or harder to maintain than what you wrote by hand?

## Why it can engage

- **Comments.** People taking a side with their own real experience
- **Saves.** The sources bundled in one place, useful to reference later
- **Shares.** Sent into a team already arguing about this
- **Quote posts.** The built-in point of this post, not a flaw to avoid

Primary metric: quote posts. Sources: LeadDev 2026-07-07 [8], METR 2025-07-10 and 2026-02-24 [9][10], 2025 DORA Report [11], Stack Overflow 2025 survey [12]. Full register in sources/sources.md.

# 09 Experiment 6. Arabic benchmarks

A 2026 quality-first Arabic leaderboard found systematic errors in benchmarks the field treats as settled, and showed a small Arabic-specialised model beating a much larger multilingual one.

> ”The Arabic benchmarks everyone cites have measurable errors. And a small Arabic-first model can beat a giant one.”

- **1.** We rank Arabic models off leaderboard averages. What if the benchmarks themselves are dirty?
- **2.** QIMMA قِمّة, built by TII, April 2026, discarded 436 of 14,163 ArabicMMLU items as flawed, a 3.1% discard rate, after quality-checking popular benchmarks [13][14].
- **3.** Scale did not decide the ranking. Fanar-1-9B and ALLaM-7B, both Arabic-specialised, outperform much larger multilingual models on specific domains [13].
- **4.** Arabic is spoken by over 400 million people and still makes up only about 0.5% of web data [15]. Native quality is a choice, not a default.
- **5.** Where has a smaller Arabic-first model beaten a bigger one for you, and which Arabic tasks still break?

> **THE QUESTION**
>
> Where has an Arabic-first model beaten a bigger frontier model for you, and which Arabic tasks still break?

## Why it can engage

- **Comments.** GCC builders naming where a native model actually won for them
- **Saves.** A reference for the next Arabic model decision
- **Shares.** Sent across Arabic-NLP and GCC developer circles
- **Real conversations.** This is the batch’s best route to one, since it sits in Muse’s own territory

Primary metric: real conversations. Sources: TII, QIMMA leaderboard and paper, April 2026 [13][14], one team’s method and finding. QCRI/HBKU, Fanar 2.0, March 2026 [15]. Lead source is slightly past 90 days, disclosed. Full register in sources/sources.md.

# 10 Measurement, a controlled learning loop

The batch runs as a learning loop, not six random posts. One primary metric per post, chosen by the post's job.

| # | JOB | PRIMARY METRIC |
|---|---|---|
| 1 | Help people choose a model | **Saves** |
| 2 | Give a workflow to try | **Saves** |
| 3 | Change how people test | **Substantive replies** |
| 4 | Prompt reflection | **Substantive replies** |
| 5 | Stake and defend an opinion | **Quote posts** |
| 6 | Open a conversation in our territory | **Real conversations** |

DEFINITIONS

**Substantive reply:** a comment that adds an opinion, a counter-point, a lived experience, or a specific question, more than one clause, not an emoji or a bare tag. **Real conversation:** a back and forth of at least three messages with a real person, or any exchange that surfaces a need, a referral, or a follow-up we act on.

Raw counts favour whichever post got amplified, so compare rates, not totals. Primary-action rate is the primary metric count divided by accounts reached, per 1,000 reached. A small post with a high save rate can and should beat a big post with a low one. `post-log.md` is the reusable template: hypothesis, format, topic, source age, channel, account, primary metric, result, quality notes, next iteration.

Secondary diagnostics, tracked but never celebrated alone: followers, likes, impressions, reach. They move for reasons unrelated to the work.

# 11 Choosing the next batch, and the source register

Rank by primary-action rate, not raw reach. Keep what earned saves, replies or conversations. Retire what earned only likes and impressions.

- Read every substantive reply and DM. Recurring questions are the next content calendar, written by the audience
- Re-run a winner in a second variant, same topic in another format or channel, before scaling it
- Watch confounds: reach gaps, which account posted, day and time, and any outside amplification

> **SAY IT PLAINLY**
>
> Six posts show direction, not proof. Treat one strong post as a hypothesis to re-run, never as a conclusion to scale on immediately.

## Condensed source register


**[1] Claude Opus 5.** Anthropic, 2026-07-24. anthropic.com/news/claude-opus-5

**[2] GPT-5.6 family.** TechCrunch, 2026-07-09. techcrunch.com

**[3] Gemini 3.6 Flash trio.** TechCrunch, 2026-07-21. techcrunch.com

**[4] Agent worktrees.** GitHub Changelog, 2026-07-30. github.blog

**[27] Worktree field report.** Leo Baniak, dev.to, 2026-07-31. dev.to

**[6] ReasonBENCH.** arXiv 2512.07795, v2 2026-05-30. arxiv.org/abs/2512.07795

**[28] AI reliance and confidence.** Tech Xplore, 2026-04-16. techxplore.com

**[22] Algorithm aversion.** Dietvorst, Simmons and Massey, 2015. doi.org/10.1037/xge0000033

**[23] Algorithm appreciation.** Logg, Minson and Moore, 2019. doi.org/10.1016/j.obhdp.2018.12.005

**[8] Code maintainability.** LeadDev, 2026-07-07, reporting GitClear research. leaddev.com

**[9][10] Developer productivity.** METR, 2025-07-10 and 2026-02-24. metr.org

**[11] The 2025 DORA Report.** DORA with Google Cloud, via Thoughtworks. thoughtworks.com

**[12] Developer survey.** Stack Overflow, 2025 AI section. survey.stackoverflow.co

**[13][14] QIMMA.** Technology Innovation Institute, 2026-04-21 and 2026-04-03. huggingface.co

**[15] Fanar 2.0.** QCRI / Hamad Bin Khalifa University, arXiv 2603.16397. arxiv.org/abs/2603.16397


Full excerpts, corrected corroborators, vendor-reported flags and study caveats live in sources/sources.md. Every cited ID maps to a URL registered through the citation-ledger script.